

EX 3 COMPARISON OF SGD WITH MOMENTUM AND ADAM OPTIMIZER (CIFAR-10)

AIM

To implement a training algorithm using stochastic gradient descent (SGD) with momentum, compare it with the Adam optimizer on the same model and dataset, and compare their convergence rate and performance.

PROCEDURE

1. Import PyTorch, torchvision and Matplotlib; select CPU/GPU.
2. Load the CIFAR-10 training set, normalise it and create a DataLoader (batch size 128).
3. Define an MLP: Flatten $\rightarrow$ Linear(3072, 256) $\rightarrow$ ReLU $\rightarrow$ Linear(256, 10).
4. Write a training function that records average loss and accuracy per epoch.
5. Train one MLP with SGD (lr = 0.01, momentum = 0.9) for 10 epochs.
6. Train an identical MLP with Adam (lr = 0.001) for 10 epochs.
7. Plot the loss and accuracy of both optimizers for comparison.

ALGORITHM

1. Start.
2. Load and normalise CIFAR-10; create mini-batches.
3. Create two identical MLP models.
4. For each optimizer (SGD + Momentum, Adam) and each epoch: forward pass, compute cross-entropy loss, zero gradients, backpropagate, update weights.
5. Accumulate the epoch loss and number of correct predictions; store the average loss and accuracy.
6. Plot loss and accuracy curves of both optimizers.
7. Compare convergence speed and final accuracy.
8. Stop.

PROGRAM

```
# Step 1: Import Necessary Libraries
import torch
import torch.nn as nn
import torch.optim as optim
import torchvision
import torchvision.transforms as transforms
import matplotlib.pyplot as plt

# Step 2: Set Device and Load Dataset
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

print("Using device:", device)

transform = transforms.Compose([
    transforms.ToTensor(),
```

```

        transforms.Normalize((0.5,),(0.5,))
    ])

trainset = torchvision.datasets.CIFAR10(
    root='./data',
    train=True,
    transform=transform,
    download=True
)

trainloader = torch.utils.data.DataLoader(
    trainset,
    batch_size=128,
    shuffle=True
)

# Step 3: Define MLP Model
class MLP(nn.Module):
    def __init__(self):
        super().__init__()

        self.flatten = nn.Flatten()

        self.net = nn.Sequential(
            nn.Linear(3 * 32 * 32, 256),
            nn.ReLU(),
            nn.Linear(256, 10)
        )

    def forward(self, x):
        x = self.flatten(x)
        return self.net(x)

# Step 4: Define Training Function
def train(model, optimizer, epochs=10):
    model.to(device)

    loss_fn = nn.CrossEntropyLoss()

    losses = []
    accuracies = []

    for epoch in range(epochs):

        total_loss = 0
        correct = 0
        total = 0

        model.train()

        for imgs, labels in trainloader:

            imgs, labels = imgs.to(device), labels.to(device)

            # Forward pass
            outputs = model(imgs)

            # Calculate loss
            loss = loss_fn(outputs, labels)

```

```

        # Backpropagation
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        # Calculate loss and accuracy
        total_loss += loss.item()

        _, predicted = outputs.max(1)

        total += labels.size(0)
        correct += (predicted == labels).sum().item()

    # Average loss and accuracy
    avg_loss = total_loss / len(trainloader)
    accuracy = 100.0 * correct / total

    losses.append(avg_loss)
    accuracies.append(accuracy)

    print(
        f"Epoch {epoch + 1}: "
        f"Loss = {avg_loss:.4f}, "
        f"Accuracy = {accuracy:.2f}%"
    )

    return losses, accuracies

# Step 5: Train with SGD + Momentum
model_sgd = MLP()

sgd = optim.SGD(
    model_sgd.parameters(),
    lr=0.01,
    momentum=0.9
)

losses_sgd, acc_sgd = train(
    model_sgd,
    sgd
)

# Step 5: Train with Adam
model_adam = MLP()

adam = optim.Adam(
    model_adam.parameters(),
    lr=0.001
)

losses_adam, acc_adam = train(
    model_adam,
    adam
)

# Step 6: Visualize Loss and Accuracy Comparison
plt.figure(figsize=(12, 5))

```

```
# Loss Comparison
plt.subplot(1, 2, 1)

plt.plot(
    losses_sgd,
    label="SGD + Momentum"
)

plt.plot(
    losses_adam,
    label="Adam"
)

plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.title("Loss Comparison on CIFAR-10 (MLP)")
plt.legend()

# Accuracy Comparison
plt.subplot(1, 2, 2)

plt.plot(
    acc_sgd,
    label="SGD + Momentum"
)

plt.plot(
    acc_adam,
    label="Adam"
)

plt.xlabel("Epoch")
plt.ylabel("Accuracy (%)")
plt.title("Accuracy Comparison on CIFAR-10 (MLP)")
plt.legend()

plt.tight_layout()
plt.show()
```


OUTPUT

```

Using device: cpu
Epoch 1: Loss = 1.7063, Accuracy = 39.99%
Epoch 2: Loss = 1.4895, Accuracy = 48.07%
Epoch 3: Loss = 1.3898, Accuracy = 51.54%
Epoch 4: Loss = 1.3138, Accuracy = 54.24%
Epoch 5: Loss = 1.2577, Accuracy = 56.19%
Epoch 6: Loss = 1.2063, Accuracy = 58.20%
Epoch 7: Loss = 1.1611, Accuracy = 59.68%
Epoch 8: Loss = 1.1194, Accuracy = 61.25%
Epoch 9: Loss = 1.0841, Accuracy = 62.50%
Epoch 10: Loss = 1.0442, Accuracy = 63.77%
Epoch 1: Loss = 1.6355, Accuracy = 42.72%
Epoch 2: Loss = 1.4316, Accuracy = 50.07%
Epoch 3: Loss = 1.3458, Accuracy = 53.13%
Epoch 4: Loss = 1.2785, Accuracy = 55.48%
Epoch 5: Loss = 1.2135, Accuracy = 57.72%
Epoch 6: Loss = 1.1594, Accuracy = 59.94%
Epoch 7: Loss = 1.1085, Accuracy = 61.70%
Epoch 8: Loss = 1.0653, Accuracy = 63.02%
Epoch 9: Loss = 1.0265, Accuracy = 64.59%
Epoch 10: Loss = 0.9686, Accuracy = 66.74%

```

Console output



Loss and accuracy comparison: SGD + Momentum vs Adam

RESULT

Both optimizers reduced the loss steadily over 10 epochs; the training accuracy of the SGD + Momentum model reached 63.77% after epoch 10. Adam started faster (42.72% accuracy in epoch 1 versus 39.99%), as seen in the plots.